


ContourRoots

Poles, zeros and roots of nonrational functions in MATLAB

Technical manual

Version 0.2.0



Flávio Luiz Cardoso-Ribeiro

<https://github.com/flavioluiz/ContourRoots>

MIT License

Abstract

ContourRoots is a MATLAB toolbox that computes the roots of scalar analytic functions, and the poles and zeros of scalar nonrational transfer functions, inside a rectangle of the complex plane. The original function is evaluated as given, without Padé approximation or modal truncation. Roots are counted with the argument principle, located by recursive subdivision and damped Newton iterations, and confirmed by local contour counts; pole-zero cancellations are resolved by counting the zeros of the numerator and denominator separately. Every result carries diagnostics that state whether the counts were reconciled, so that incomplete searches are reported rather than hidden. For scalar retarded time-delay systems, the toolbox also computes every critical delay and crossing direction directly and gives a bound-based count of unstable roots. This manual describes the mathematical background, the algorithm and its numerical safeguards, the interface, and three families of case studies: time-delay systems and the limits of Padé approximations, distributed-parameter systems, and a beam coupled to an oscillator. Proofs are collected in the appendices.

Contents

1	Introduction	3
1.1	The problem	3
1.2	What the toolbox provides	3
1.3	What the toolbox does not claim	4
1.4	How to read this manual	4
2	Quick start	5
2.1	Installation	5
2.2	Roots of a characteristic function	5
2.3	Poles and zeros	5
2.4	Time-delay stability	6
3	Mathematical background	7
3.1	Analytic and meromorphic functions	7
3.2	The argument principle	7
3.3	Cancellation of poles and zeros	8
3.4	Modes and transfer poles	9
4	Algorithm and numerical safeguards	10
4.1	Contour counts	10
4.2	Newton's method	10
4.3	Local confirmation	11
4.4	Recursive subdivision	11
4.5	Cancellation	11
4.6	Exploratory mode	12
4.7	The completeness contract	12
4.8	Cost	12
4.9	Related methods and software	12
5	Model input and the analyticity contract	13
5.1	Input forms	13
5.2	The symbolic analyticity check	13
5.3	What a declaration of analyticity means	13
5.4	Singular points	14
5.5	Scaling	14
6	Interface reference	15
6.1	Functions	15
6.2	Options	15
6.3	Diagnostics	16

7	Time-delay systems and the limits of Padé approximations	17
7.1	Formulation	17
7.2	Structure of the spectrum	18
7.3	Roots for a given delay	19
7.4	Critical delays	19
7.5	Direction of the crossings	20
7.6	Stability map	20
7.7	The Padé approximation	21
7.8	What Padé preserves	23
7.9	What Padé loses	24
7.9.1	A non-conservative delay margin	24
7.9.2	Shifted and missing branches	24
7.9.3	Stability windows erased, shifted and invented	26
7.10	Guidelines	28
7.11	Limitations	28
8	Distributed-parameter systems	29
8.1	Removing apparent branch points	29
8.2	Heat equation: the sensor position	29
8.3	String and duct	30
8.4	Kelvin–Voigt beam: an accumulation point	31
8.5	Results	32
9	A beam coupled to an oscillator	33
9.1	Model	33
9.2	The bare-beam receptance	34
9.3	Characteristic function and transfer functions	34
9.4	Implementation	34
9.5	Reference case and parameter studies	35
9.6	Independent validation	36
10	Validation and limitations	38
10.1	Test suite	38
10.2	Independent references in the studies	38
10.3	Limitations	38
A	Proofs	40
B	Reproducibility	41

Chapter 1

Introduction

1.1 The problem

Many questions about linear dynamical systems reduce to finding the values of a complex variable s where a scalar function vanishes. The modes of a system are the roots of its characteristic function $\Delta(s)$; the poles and zeros of a transfer function $G(s) = N(s)/D(s)$ are the roots of D and N that survive cancellation. When Δ , N and D are polynomials, the problem is solved by computing matrix eigenvalues, as MATLAB's `roots`, `pole` and `zero` do.

Two large classes of models are not rational:

- systems with *time delays*, whose characteristic functions contain terms such as e^{-sT} [3, 9, 10, 14];
- *distributed-parameter* systems, described by partial differential equations, whose transfer functions contain $\sinh$, $\cosh$, $\sqrt{s}$ and similar functions [7].

Their characteristic functions typically have infinitely many roots, and no finite matrix contains all of them. A common practice is to replace the nonrational terms by rational approximations, for example Padé approximations of e^{-sT} or truncated modal expansions. The approximation makes the finite-dimensional toolbox available again, but it changes the question: one then computes the roots of a different function. Chapter 7 shows, on simple examples, that the answer can change qualitatively.

ContourRoots answers the original question in a bounded region: *which roots does the function have inside this rectangle?* It evaluates the function as it is given, and uses the argument principle to count how many roots each part of the rectangle contains, which makes it possible to tell whether the search was complete.

1.2 What the toolbox provides

- `croots`, `cpoles`, `czeros` and `cpzmap`, a MATLAB-style interface for roots, poles, zeros and pole-zero maps of scalar functions given as function handles, coefficient vectors, numerator/denominator pairs (`ndpair`), symbolic expressions or SISO LTI models.
- Detection of pole-zero cancellations, including partial ones.
- Diagnostics for every search: a status, the multiplicity and residual of each location, cancelled locations and unresolved regions.
- For scalar retarded time-delay systems $D(s) + N(s)e^{-sT} = 0$: all critical delays and crossing directions up to a given delay, without approximation, and a count of unstable roots in a region that provably contains all of them.
- Examples, reproducible studies and tests.

1.3 What the toolbox does not claim

The word *exact* is used in this manual only in one sense: the original function is evaluated, without rational approximation or truncation. The computed roots are floating-point numbers. A search that ends with the status `numerically_complete` has passed careful numerical checks under the analyticity assumption, but it is not a proof in the sense of interval arithmetic (Section 4.7). The toolbox does not treat matrix-valued characteristic functions, does not enumerate infinite spectra, and does not discover singularities hidden inside opaque function handles. No claim of superiority over other root finders is made; related methods and software are discussed in Section 4.9.

1.4 How to read this manual

Chapter 2 is a short tour of the interface. Chapters 3–5 give the mathematical background, the algorithm, and the precise meaning of each input form and of the analyticity assumption. Chapter 6 is a compact reference. Chapters 7–9 are case studies that can be read independently: time-delay systems and the limits of Padé approximations; distributed-parameter systems; and a beam coupled to an oscillator. Chapter 10 summarizes the validation, and the appendices contain proofs and reproducibility information.

The Markdown documentation distributed with the toolbox covers the same material in a tutorial style, with runnable code for every step.

Chapter 2

Quick start

2.1 Installation

ContourRoots is plain MATLAB code; the core needs no additional toolbox. Download ContourRoots.mltbx from the release page and open it to install it as an Add-On, or download ContourRoots.zip, unzip it and run `setup_contourroots`, which adds the toolbox to the path for the current session. From the Command Window:

```
url = "https://github.com/flavioluiz/ContourRoots/releases/latest/download/
      ContourRoots.zip";
unzip(websave("ContourRoots.zip", url), "ContourRoots");
run("ContourRoots/setup_contourroots.m")
```

2.2 Roots of a characteristic function

The characteristic function $F(s) = s^2 + s + 1 + (2s + 3)e^{-s}$ has infinitely many roots. Those in $-8 < \operatorname{Re} s < 2$, $-20 < \operatorname{Im} s < 20$ are obtained with

```
F = @(s) s.^2 + s + 1 + (2*s + 3).*exp(-s);
[r, info] = croots(F, [-8 2 -20 20], 'AssumeAnalytic', true);
```

The vector `r` contains seven roots, ordered by decreasing real part; the first two, $0.36201 \pm 1.8195i$, lie in the right half-plane. `info.status` is `'numerically_complete'`. The option `'AssumeAnalytic', true` is the user's statement that F is analytic on a neighborhood of the closed rectangle; it is justified here because F is built from polynomials and the exponential. Without it the search is exploratory (Section 4.6).

2.3 Poles and zeros

A transfer function is given as a pair of analytic functions:

```
G = ndpair(@(s) sinh(s/2), @(s) sinh(s));
p = cpoles(G, [-1 1 -10 10], 'AssumeAnalytic', true); % +-i*pi, +-3i*pi
z = czeros(G, [-1 1 -10 10], 'AssumeAnalytic', true); % empty
cpzmap(G, [-1 1 -10 10], 'AssumeAnalytic', true)
```

The denominator vanishes at $i\pi k$ for every integer k ; for even k the numerator vanishes as well and the candidate cancels. Figure 2.1 shows both results.




Figure 2.1: Left: roots of $s^2 + s + 1 + (2s + 3)e^{-s}$ in the rectangle $[-8, 2] \times [-20, 20]$ (dashed); the shaded area is the right half-plane. Right: `cpzmap` of $\sinh(s/2)/\sinh(s)$; grey dots are cancelled candidates.

2.4 Time-delay stability

For $D(s) + N(s)e^{-sT} = 0$ with coefficient vectors $\mathbf{N}$ and $\mathbf{D}$:

```
crit = critical_delays(2, [1 1], 10);      % s + 1 + 2 e^{-sT} = 0
Z = unstable_root_count(2, [1 1], 1.5);    % 2 unstable roots
```

The first critical delay is $2\pi/(3\sqrt{3}) \approx 1.2092$ (Chapter 7).

Chapter 3

Mathematical background

This chapter collects the facts from complex analysis on which the method relies; see Ahlfors [1] for proofs of the classical results.

3.1 Analytic and meromorphic functions

A function F of the complex variable $s = x + iy$ is *analytic* (holomorphic) on an open set U if it has a complex derivative at every point of U . It is *entire* if it is analytic on the whole plane. Polynomials, e^s , $\sin s$, $\cos s$, $\sinh s$, $\cosh s$, and all sums, products and compositions of entire functions are entire. A function is *meromorphic* on U if it is analytic on U except at isolated poles.

Two properties of analytic functions are essential:

1. The zeros of an analytic function that is not identically zero are isolated, and each has a finite *multiplicity* $m \geq 1$: $F(s) = (s - s_0)^m g(s)$ with g analytic and $g(s_0) \neq 0$.
2. If F is analytic on a neighborhood of a closed bounded region and not identically zero, that region contains finitely many zeros. The hypothesis on the closure matters: zeros may accumulate at a singular point on the boundary of a region (see Section 8.4).

Neither holds for functions that are not analytic: $|s|^2 - 1$ vanishes on a whole circle. Functions with branch points (such as $\sqrt{s}$ or $\log s$) are analytic only on a domain that excludes a cut.

Some expressions that contain square roots are nevertheless entire. For instance $\cosh \sqrt{q} = \sum_k q^k / (2k)!$ and $\sinh(\sqrt{q}) / \sqrt{q} = \sum_k q^k / (2k + 1)!$ contain only integer powers of q : the choice of the branch of $\sqrt{q}$ does not change their value. Such *even functions of a root* appear in the transfer functions of Chapters 8 and 9.

3.2 The argument principle

Theorem 3.1 (Argument principle). *Let C be a simple closed piecewise-smooth curve, oriented counterclockwise, and let F be meromorphic on a neighborhood of C and its interior, with no zero or pole on C . Then*

$$\frac{1}{2\pi i} \oint_C \frac{F'(s)}{F(s)} ds = Z - P = \frac{1}{2\pi} \Delta_C \arg F(s), \quad (3.1)$$

where Z and P are the numbers of zeros and poles inside C , counted with multiplicity, and $\Delta_C \arg F$ is the total change of a continuous branch of $\arg F(s)$ as s travels once around C .

The right-hand form is the one used numerically: the change of argument is the sum of the phase increments of F between consecutive samples of C , each taken in $(-\pi, \pi]$. It equals 2π times the *winding number* of the image curve $F(C)$ around the origin (Figure 3.1).




Figure 3.1: The argument principle for $F(s) = s^2 + s + 1 + (2s + 3)e^{-s}$ on the rectangle $[-2, 1] \times [-3, 9]$. (a) Four roots lie inside C . (b) The image curve $F(C)$, drawn with radius $\log(1 + |F|)$, winds four times around the origin. (c) The unwrapped argument of F along C grows by $4 \times 2\pi$.

Three consequences shape the design of ContourRoots.

Analytic functions. If F is analytic, $P = 0$ and the integral counts the zeros. This is the case for characteristic functions such as $D(s) + N(s)e^{-sT}$.

Quotients. For a transfer function $G = N/D$ with N and D analytic, the integral of G'/G counts $Z_N - Z_D$ after cancellation. A count of zero does not mean that the region is empty. This is why ContourRoots asks for N and D separately (`ndpair`) and counts the zeros of each.

Multiplicity and the first moment. If F is analytic with zeros $s_1, \dots, s_Z$ (repeated according to multiplicity) inside C ,

$$\frac{1}{2\pi i} \oint_C s \frac{F'(s)}{F(s)} ds = \sum_{j=1}^Z s_j. \quad (3.2)$$

Divided by Z it gives the centroid of the zeros; when $Z = 1$ it gives the zero itself. Higher moments lead to the classical methods of Delves and Lyness [8] and Kravanja and Van Barel [12], which compute all zeros inside C from a small eigenvalue problem. ContourRoots uses only the first moment, as a starting point for Newton's method.

3.3 Cancellation of poles and zeros

Let N and D be analytic near s_0 , with zeros of orders $m_N \geq 0$ and $m_D \geq 0$ at s_0 . Then $G = N/D$ has at s_0

- a pole of order $m_D - m_N$ if $m_D > m_N$;
- a zero of order $m_N - m_D$ if $m_N > m_D$;
- neither (a removable point, after cancellation) if $m_N = m_D$.

Counting the zeros of N and of D separately in a small disk around each candidate determines m_N and m_D , and therefore the type and order of the point. This is the cancellation test of Section 4.5.

3.4 Modes and transfer poles

For a linear system, the roots of the characteristic function are the *modes*; they decide internal stability. A transfer function from one input to one output may not contain every mode: a mode that the input cannot excite, or that the output cannot observe, appears as a common zero of N and D and cancels. Both computations are available: `croots` on the characteristic function returns the modes, and `cpoles` on N/D returns the transfer poles. Examples of hidden modes appear in Sections 8.2 and 9.6.

Chapter 4

Algorithm and numerical safeguards

The solver, `complex_spectrum`, isolates the zeros of an analytic function f in a rectangle $\Omega = [x_0, x_1] \times [y_0, y_1]$. In zero mode $f = N$ (or the given function); in pole mode $f = D$. This chapter describes each component and the conditions under which a result is declared complete.

4.1 Contour counts

The boundary of a rectangle B is sampled at n equally spaced points per edge, $4n$ points in total, traversed counterclockwise. With $v_j = f(z_j)$, the phase and log-modulus increments are

$$\delta_j = \text{Arg} \frac{v_{j+1}}{v_j} \in (-\pi, \pi], \quad \mu_j = \log |v_{j+1}| - \log |v_j|,$$

and the discrete winding number is $w = \frac{1}{2\pi} \sum_j \delta_j$. The number of samples starts at $n = n_0$ (`ContourPoints`, default 32) and is doubled up to `ContourRefinements` times (default 7, i.e. up to 4096 points per edge). A count is *accepted* when

1. the rounded winding number $\text{round}(w)$ is the same at three consecutive sampling levels,
2. at the last two levels $|w - \text{round}(w)| < 10^{-7}$, and
3. at the last two levels every increment satisfies $|\delta_j| < \pi/3$ and $|\mu_j| < 2$.

Condition 3 is the practical form of the requirement that the argument of f vary slowly between samples; without it, a fast phase rotation could be aliased and a whole turn lost. If any sample is zero, infinite or NaN (a zero on the boundary, or overflow), or if the conditions are not met at the finest level, the count fails and the rectangle is *unresolved*.

The same samples give the first-moment estimate (3.2), computed as $\sum_j \frac{z_j + z_{j+1}}{2} (\mu_j + i\delta_j) / (2\pi i Z)$, used as a Newton starting point.

4.2 Newton's method

From a starting point z inside a cell B with count m , the solver applies the modified Newton iteration

$$z \leftarrow z - \alpha m \frac{f(z)}{f'(z)},$$

which converges quadratically to a zero of multiplicity m . Safeguards:

- the step is limited to half the larger side of B ;
- the damping factor α is halved (at most 12 times) until $|f|$ does not increase; if it cannot be decreased, the start is abandoned;

- the iteration stops if it leaves B , and converges when the step is below $\text{RootTolerance} \cdot (1 + |z|)$.

The derivative is the user-supplied `Derivative`, the exact derivative for symbolic input and coefficient vectors, or otherwise the four-point complex difference

$$f'(z) \approx \frac{f(z+h) - f(z-h) - i[f(z+ih) - f(z-ih)]}{4h}, \quad h = 10^{-4} \max(1, |z|),$$

whose error is $O(h^4)$ for analytic f (the h^2 terms of the real and imaginary central differences cancel).

4.3 Local confirmation

A Newton limit z^* is accepted only if a contour count on the small square of half-width $\rho = \text{RootTolerance} \cdot (1 + |z^*|)$ centered at z^* equals the count m of the cell. This confirms both that z^* is a zero and that no other zero of the cell was missed; m becomes the multiplicity of z^* , and $\sqrt{2}\rho$ is reported as `locationRadius`. Consequently, zeros closer than about ρ are reported as one location whose multiplicity is the size of the cluster.

4.4 Recursive subdivision

Algorithm 1 — recursive isolation of the zeros of f in a cell B with count m

1. If $m = 0$, return.
2. For each start in (first-moment center of B , geometric center of B , user seeds inside B): run Newton (Section 4.2); if the limit is confirmed with count m (Section 4.3), record it and return.
3. If the depth exceeds `MaxDepth` or the number of visited cells exceeds `MaxCells`, mark B unresolved and return.
4. Split B across its longer side at the fraction 0.5 of the side. Count the two halves. If both counts are accepted and add up to m , recurse into both halves and return.
5. Otherwise retry with the split fractions 0.438... and 0.573..., then with the other side. If no split conserves the count, mark B unresolved.

The *conservation check* of step 4 detects zeros lying on the split line (the counts of the halves then fail or do not add up) and moves the line. The irrational-looking fractions make it unlikely that a second attempt hits a zero again. Unresolved cells are never discarded: they are returned in `info.unresolvedBoxes`.

4.5 Cancellation

In pole mode with a pair N/D , the zeros of D are isolated with Algorithm 1. For each location z^* with multiplicity m_D , the zeros of N are counted on the same small square, giving m_N . The reported multiplicity is $\max(m_D - m_N, 0)$; if it is zero, z^* is moved to `cancelledLocations`, and the removed order $\min(m_D, m_N)$ is recorded. Zero mode is symmetric. If the count of the other factor fails, the location is kept and flagged in `cancellationUncertain`, and the result is not complete. The test has the resolution of the small square: a pole and a zero closer than ρ are indistinguishable from an exact cancellation.

4.6 Exploratory mode

When the function is not known to be analytic (an opaque handle without `AssumeAnalytic`, or a single quotient handle), contour counts may measure $Z - P$ instead of Z and cannot be trusted. The solver then starts Newton from a `GridSize` grid of points and from the seeds, and keeps the limits confirmed by a positive local count. The result has status `exploratory` and is never declared complete.

4.7 The completeness contract

A search is `numerically_complete` when (i) the count on $\partial\Omega$ was accepted, (ii) every cell was resolved, (iii) the multiplicities of the located zeros add up to that count, and (iv) every cancellation test was resolved. Under the analyticity assumption, and if the accepted counts are correct, all zeros in Ω are then reported, each within its `locationRadius` as a cluster of the stated multiplicity.

What this is not

It is not a proof. The counts are computed from finitely many floating-point samples; a function that oscillates between samples faster than the safeguards of Section 4.1 can detect could in principle fool them. Analyticity of a handle is the user's declaration. Residuals depend on the scaling of the searched factor and are not bounds on the location error. `info.certified` is always false. Rigorous enclosures would require interval arithmetic, which is outside the scope of this version.

4.8 Cost

The cost is dominated by function evaluations on contours: $4n$ per count, with n doubling until acceptance, for the outer rectangle, for every cell of the subdivision and for every local confirmation. Vectorized handles are evaluated on all samples of a contour at once. Tall rectangles over oscillatory functions (e.g. e^{-sT} with large $T|\operatorname{Im} s|$) need more samples per edge; smaller rectangles or larger `ContourRefinements` help. Seeds from a previous parameter value reduce the number of subdivisions in parameter sweeps.

4.9 Related methods and software

Contour-integral root finders go back to Delves and Lyness [8]; see Kravanja and Van Barel [12] for a thorough treatment including moment-based eigenvalue formulations. `cxroots` [16] implements such a method in Python with subdivision. GRPF [11] analyzes the phase of a single meromorphic function on an adaptive triangulation and locates both zeros and poles. For quasi-polynomials and time-delay systems, QPmR [18] maps the zero-level curves of the real and imaginary parts, and spectral discretization methods [4, 5, 14] approximate the rightmost roots of matrix delay equations. ContourRoots combines well-known ingredients; its emphasis is on a MATLAB-style interface, explicit cancellation for N/D pairs, and diagnostics that make incomplete searches visible.

Chapter 5

Model input and the analyticity contract

The algorithm of Chapter 4 is only as good as its assumption that the searched factor is analytic on a neighborhood of the closed rectangle. This chapter lists the accepted input forms and states, for each, how that assumption is established.

5.1 Input forms

Input	Treatment
Coefficient vector	Polynomial in descending powers, exact derivative; analytic by construction. <code>croots(c,region)</code> behaves like <code>roots(c)</code> restricted to the rectangle.
Function handle	Evaluated as given. Analytic only if the user passes <code>AssumeAnalytic</code> ; otherwise exploratory. Vectorized handles are evaluated on whole contours; scalar-only handles are evaluated point by point.
<code>ndpair(N,D)</code>	Each factor is a handle, a coefficient vector or a symbolic expression. Polynomial and checked symbolic factors are analytic automatically; handles need <code>AssumeAnalytic</code> , which then applies to <i>both</i> factors.
Symbolic (<code>sym</code> , <code>symfun</code>)	Split with <code>numden</code> ; each factor is checked by a conservative grammar (below) and differentiated exactly.
<code>tf</code> , <code>zpk</code> , <code>ss</code>	Continuous-time SISO only. Rational models and <code>tf</code> models with input, output or transport delays become an analytic pair (delay in the numerator). Models with internal delays are evaluated with <code>evalfr</code> in exploratory mode.

5.2 The symbolic analyticity check

A symbolic factor is accepted as entire if it is a number, the variable, or is built from entire factors by sums, products, exp, sin, cos, sinh, cosh and powers with constant nonnegative integer exponents. The check is *sufficient, not necessary*: $\cosh \sqrt{s}$ is entire but is rejected, with the error `complex_spectrum:AnalyticContract`. The remedy is to provide the factor as a handle that evaluates the entire function directly (with a series near removable points, Chapter 8) and to declare `AssumeAnalytic`. All parameters other than the complex variable must be substituted numerically before the call.

5.3 What a declaration of analyticity means

`'AssumeAnalytic',true` asserts that the searched factor(s) are analytic on an open set containing the closed rectangle. It is a statement about mathematics, not about the code, and must be checked by the user:

- removable singularities (e.g. $\sinh(s)/s$ at 0) must be removed in the implementation, otherwise the evaluation returns `NaN`;
- roots such as $\sqrt{s}$ or $s^{1/4}$ are acceptable only inside combinations invariant under the change of branch;
- for a *single* handle in pole mode, the declaration concerns $1/G$: it asserts that G has no zeros in the rectangle. For transfer functions with zeros, use `ndpair`.

A false declaration makes the counts meaningless and can produce wrong results without any warning.

5.4 Singular points

Some models have points where the function is not meromorphic: branch points of genuinely multivalued functions, or accumulation points of poles (Section 8.4). No finite contour computation can resolve a neighborhood of such a point. The option `Singularities` declares them; a rectangle that contains one is rejected with the error `complex_spectrum:SingularityInRegion`. The solver does not discover such points in opaque handles.

5.5 Scaling

The argument principle is invariant under multiplication of f by a nonzero constant, and the regression tests check that roots do not change when f is scaled by $10^{\pm 100}$. Overflow is a different matter: a term such as e^{-sT} grows like $e^{T|\operatorname{Re} s|}$ in the left half-plane, and samples that overflow make a count fail. Rectangles extending far to the left may need to be split, or the function rescaled by an analytic nonvanishing factor.

Chapter 6

Interface reference

The complete reference, with examples for every function, is in `docs/api/` and in the MATLAB help (`help croots`). This chapter summarizes it.

6.1 Functions

Syntax	Purpose
<code>[r,info] = croots(F,region,...)</code>	roots of F in the rectangle
<code>[p,info] = cpoles(G,region,...)</code>	poles of G after cancellation
<code>[z,info] = czeros(G,region,...)</code>	zeros of G after cancellation
<code>cpzmap(G,region,...)</code>	pole-zero map; <code>[p,z,infoP,infoZ]</code> without plot
<code>G = ndpair(N,D)</code>	transfer function from two analytic factors
<code>[x,info] = complex_spectrum(...)</code>	core solver, 'Mode' 'zeros'/'poles'
<code>crit = critical_delays(N,D,Tmax)</code>	crossings of $D + Ne^{-sT}$, $0 \leq T \leq T_{\max}$
<code>Z = unstable_root_count(N,D,T)</code>	roots with $\operatorname{Re} s \geq 0$ (bound-based count)
<code>R = rhp_root_bound(N,D)</code>	radius containing all such roots, all T
<code>delay_roots, delay_root_count</code>	grid/Newton solver and count for one delay
<code>pade_delay, pade_characteristic</code>	diagonal Padé coefficients and polynomial
<code>pade_critical_delays</code>	critical delays predicted by Padé models
<code>setup_contourroots, contourroots</code>	path setup; overview and version

The region is always `[xmin xmax ymin ymax]`. When `info` is not requested and a search of `croots`, `cpoles`, `czeros` or `cpzmap` is not complete, a warning (`ContourRoots:Exploratory` or `ContourRoots:Incomplete`) is issued; `'Warn',false` silences it.

6.2 Options

Option	Default	Meaning
<code>AssumeAnalytic</code>	false	declare the searched factor(s) analytic
<code>SeedPoints</code>	<code>[]</code>	extra Newton starts
<code>Derivative</code>	<code>[]</code>	F' for a single handle
<code>Singularities</code>	<code>[]</code>	points that must not lie in the region
<code>RootTolerance</code>	1e-7	relative half-width of confirmation boxes
<code>ContourPoints</code>	32	initial samples per edge
<code>ContourRefinements</code>	7	maximum doublings of the sampling
<code>MaxDepth</code>	24	maximum subdivision depth
<code>MaxCells</code>	5000	maximum number of cells
<code>MaxIterations</code>	80	Newton iterations per start
<code>GridSize</code>	[25 41]	start grid in exploratory mode
<code>Plot, Display</code>	false	figure or table of results
<code>Warn</code>	true	incompleteness warning (wrappers only)

6.3 Diagnostics

Field	Meaning
<code>status</code>	<code>numerically_complete</code> , <code>unresolved</code> or <code>exploratory</code>
<code>complete</code>	conditions (i)–(iv) of Section 4.7 hold
<code>certified</code>	always false
<code>count</code>	total multiplicity returned, after cancellation
<code>targetCountBeforeCancellations</code>	outer count of the searched factor
<code>multiplicity</code>	multiplicity (cluster size) of each location
<code>locationRadius</code>	resolution of each location (not an error bound)
<code>residuals</code>	$ f $ of the searched factor at each location
<code>cancelledLocations,</code> <code>cancellationOrders</code>	removed candidates and orders
<code>cancellationUncertain</code>	cancellation test unresolved at a location
<code>unresolvedBoxes</code>	cells that could not be resolved
<code>contourResolved, cellsVisited</code>	outer count accepted; cells used
<code>exploratory, region, mode,</code> <code>analyticSource</code>	context of the search

Chapter 7

Time-delay systems and the limits of Padé approximations

This chapter studies the scalar characteristic equation

$$F(s, T) = D(s) + N(s) e^{-sT} = 0, \quad T \geq 0, \quad (7.1)$$

with real polynomials D of degree m and N of degree p . It shows that the stability of such systems can be decided for every delay without approximating the delay, and then examines precisely what a Padé approximation of e^{-sT} preserves and what it loses.

7.1 Formulation

The simplest example is the delay-differential equation

$$\dot{x}(t) = -a x(t) - b x(t - T). \quad (7.2)$$

Its solution from $t = 0$ is determined only by a whole *history* $x(\theta) = \varphi(\theta)$, $\theta \in [-T, 0]$: on $[0, T]$ the delayed term is known and (7.2) is an ordinary differential equation; its solution determines the delayed term on $[T, 2T]$, and so on (the *method of steps* [3]). The state is a function, and the system is infinite-dimensional. Exponential solutions e^{st} exist when $s + a + be^{-sT} = 0$. More generally, the loop of Figure 7.1, a rational plant N/D with delayed unit feedback, has the characteristic equation (7.1).




Figure 7.1: Loop with a delay in the feedback path. The loop transfer function is $L(s) = \frac{N(s)}{D(s)} e^{-sT}$; $1 + L = 0$ is equivalent to (7.1). The delay has unit gain at every frequency and unbounded phase $-\omega T$.

Definition 7.1. The system is *retarded* if $p < m$ and *neutral* if $p = m$.

Throughout this chapter $p < m$. Three cases illustrate the typical behaviors (Table 7.1).

Table 7.1: Case studies; all are stable at $T = 0$.

Case	Equation	Model	Behavior as T grows
P1	$s + 1 + 2e^{-sT} = 0$	$\dot{x} = -x(t) - 2x(t - T)$	loses stability at $T \approx 1.209$
P2	$s + 2 + e^{-sT} = 0$	$\dot{x} = -2x(t) - x(t - T)$	stable for every $T \geq 0$
P3	$s^2 + 0.8s + 4 + 3e^{-sT} = 0$	$\ddot{x} + 0.8\dot{x} + 4x = -3x(t - T)$	stability switches

7.2 Structure of the spectrum

$F(\cdot, T)$ is entire and, for $T > 0$, has infinitely many roots. For retarded systems they form chains that move to the left: balancing $|D(s)| \sim |s|^m$ with $|N(s)e^{-sT}| \sim |s|^p e^{-T \operatorname{Re} s}$ gives

$$\operatorname{Re} s \approx -\frac{m-p}{T} \ln |s|, \quad |s| \rightarrow \infty.$$

Hence every half-plane $\operatorname{Re} s > c$ contains finitely many roots, the *spectral abscissa* $\alpha(T) = \max\{\operatorname{Re} s : F(s, T) = 0\}$ is well defined, and the zero solution of a linear retarded system is exponentially stable if and only if $\alpha(T) < 0$ [10, 14]. Figure 7.2 shows the spectra of the three cases together with the poles of Padé models (Section 7.7).




Figure 7.2: Roots of (7.1) ($\times$) and poles of diagonal Padé models ($\circ$) at $T = 1$ (P1, P2) and $T = 0.8$ (P3). The exact chains continue to the left; the rational spectra are finite.

Proposition 7.2 (A region containing all unstable roots). *Let $p < m$, D monic, $D(s) = s^m + \sum_{k < m} d_k s^k$, $N(s) = \sum_{k \leq p} n_k s^k$, and let $R > 0$ be the positive root of*

$$r^m - \sum_{k < m} |d_k| r^k - \sum_{k \leq p} |n_k| r^k = 0.$$

Then for every $T \geq 0$, every root of (7.1) with $\operatorname{Re} s \geq 0$ satisfies $|s| \leq R$.

The proof is in Appendix A. For P1, $|s + 1| \leq 2$ gives $R = 3$; for P3, $R \approx 3.076$. The function `rhp_root_bound` computes R , and `unstable_root_count` applies the argument principle to the rectangle $[0, R + 1] \times [-(R + 1), R + 1]$: the number $Z(T)$ of unstable roots is obtained without guessing a search window and without approximation (the count itself is a floating-point computation with adaptive sampling; since the left edge is the imaginary axis, it is inconclusive, and returns `NaN`, when roots lie on or extremely close to the axis — near critical delays, for persistent imaginary roots, or for large delays, where many roots approach the axis. $Z(T) = 0$ means stability only when no root lies on the imaginary axis).

7.3 Roots for a given delay

The study of this chapter locates roots with the delay-specific solver `delay_roots` (damped Newton from a grid, audited by an argument-principle count); `croots` gives the same roots with the subdivision algorithm of Chapter 4. Table 7.2 shows that the number of roots found agrees with the independent count and that the residuals are near machine precision.

Table 7.2: Solver diagnostics at the nominal delay.

Case	Roots	Count	Max residual	Abscissa
P1	14	14	1.13e-13	-0.09248
P2	14	14	4.46e-14	-0.86098
P3	12	12	2.18e-12	0.22028

7.4 Critical delays

As T varies the roots move continuously, and new roots appear only from $\operatorname{Re} s = -\infty$. Stability can therefore change only when a root crosses the imaginary axis. At a crossing $s = i\omega$, $\omega > 0$,

$$e^{-i\omega T} = -\frac{D(i\omega)}{N(i\omega)}. \quad (7.3)$$

This complex equation splits into two real equations with very different roles.

Modulus: where. The left side has unit modulus for every T , so the crossing frequencies do not depend on the delay and solve

$$G(\omega) = |D(i\omega)|^2 - |N(i\omega)|^2 = 0, \quad (7.4)$$

an even polynomial in ω , i.e. a polynomial of degree m in $x = \omega^2$. Its positive roots give a *finite* set of frequencies.

Phase: when. For each frequency ω_i , with $\theta_i = \operatorname{mod}(-\arg[-D(i\omega_i)/N(i\omega_i)], 2\pi)$, the phase condition gives *infinitely many* critical delays

$$T_{i,k} = \frac{\theta_i + 2\pi k}{\omega_i}, \quad k = 0, 1, 2, \dots \quad (7.5)$$

This elimination is the scalar form of the direct method [14, 15, 19]. `critical_delays` computes all $T_{i,k}$ in an interval and refines each pair (ω, T) by a two-variable Newton correction.

Example 7.3 (P1). $G(\omega) = 1 + \omega^2 - 4$ gives $\omega_c = \sqrt{3}$. Since $-D(i\sqrt{3})/N = e^{-2\pi i/3}$, $\theta = 2\pi/3$ and $T_k = (2\pi/3 + 2\pi k)/\sqrt{3}$: $T_0 = 2\pi/(3\sqrt{3}) \approx 1.2092$, $T_1 \approx 4.8368$, $T_2 \approx 8.4644$.

Example 7.4 (P3). $G(\omega) = \omega^4 - 7.36\omega^2 + 7$ has the roots $\omega_1 \approx 2.4976$ and $\omega_2 \approx 1.0593$; each generates its own sequence of critical delays, and the two sequences interleave.

For P2, $G(\omega) = \omega^2 + 3 > 0$: there is no crossing, and since P2 is stable at $T = 0$ it is stable for every delay.

7.5 Direction of the crossings

Differentiating $F(s(T), T) = 0$ gives

$$\frac{ds}{dT} = -\frac{F_T}{F_s} = \frac{s e^{-sT} N(s)}{D'(s) + e^{-sT} [N'(s) - TN(s)]}. \quad (7.6)$$

The crossing is destabilizing if $\operatorname{Re}(ds/dT) > 0$ at $s = i\omega$. The following classical result [6] makes the direction a property of the frequency alone.

Proposition 7.5 (Direction depends only on the frequency). *At a simple crossing $s = i\omega$, $\omega > 0$, $\operatorname{sign} \operatorname{Re} \frac{ds}{dT} = \operatorname{sign} G'(\omega)$.*

All branches k of a frequency therefore cross in the same direction. In P1, $G'(\sqrt{3}) > 0$: every T_k adds an unstable pair. In P3, crossings at ω_1 always destabilize and crossings at ω_2 always stabilize (Table 7.3, Figure 7.3).

Table 7.3: Imaginary-axis crossings in the studied delay range.

Case	T_c	ω_c	branch	$\operatorname{Re}(ds/dT)$	direction
P1	1.20919958	1.73205081	0	+3.237e-01	destabilizing
P2			no crossing		
P3	0.29182687	2.49756112	0	+1.025e+00	destabilizing
P3	2.69528299	1.05933396	0	-6.200e-02	stabilizing
P3	2.80755521	2.49756112	1	+2.239e-01	destabilizing




Figure 7.3: Frequencies and directions of the exact crossings (up: into the right half-plane; down: back to the left half-plane).

7.6 Stability map

Each simple crossing moves one conjugate pair, so the number of roots in the right half-plane is the step function

$$Z(T) = Z(0) + 2 \sum_{T_{i,k} < T} \operatorname{sign} G'(\omega_i), \quad (7.7)$$

where $Z(0)$ is the number of unstable roots of $D + N$ (the roots created when T becomes positive come from $\operatorname{Re} s = -\infty$). For P3 on $[0, 10]$:

$$Z(T) = \begin{cases} 0, & 0 \leq T < 0.2918, \\ 2, & 0.2918 < T < 2.6953, \\ 0, & 2.6953 < T < 2.8076 \quad (\text{stability window}), \\ 2, 4, 6, 4, & \text{after } 2.8076, 5.3233, 7.8390, 8.6265. \end{cases}$$

The window has width 0.112: a sweep with step 0.5 misses it entirely. Formula (7.7) was checked independently with `unstable_root_count` at 80 delays, with full agreement (black curve and circles in Figure 7.12). Figure 7.4 shows the spectral abscissa computed by a sweep, refined near the window.




Figure 7.4: Spectral abscissa versus delay, exact and with Padé models of orders 1, 4 and 8; diamonds are the exact critical delays of (7.5). In P3 the exact curve touches zero at $T \approx 2.695$ and crosses again at $T \approx 2.808$; the $[1/1]$ model misses the window and the $[4/4]$ model widens it.

7.7 The Padé approximation

The diagonal Padé approximation of e^{-z} of order n is the rational function $P_n(z)/Q_n(z)$ whose Taylor series agrees with that of e^{-z} up to z^{2n} [2]:

$$e^{-z} - \frac{P_n(z)}{Q_n(z)} = O(z^{2n+1}), \quad Q_n(z) = \sum_{k=0}^n \frac{(2n-k)!n!}{(2n)!k!(n-k)!} z^k, \quad P_n(z) = Q_n(-z). \quad (7.8)$$

With $z = sT$ the characteristic equation becomes the polynomial $D(s)Q_n(sT) + N(s)P_n(sT)$ of degree $m + n$. Three structural properties determine its behavior.

(i) Exact gain. Since $P_n(iy) = \overline{Q_n(iy)}$,

$$\left| \frac{P_n(iy)}{Q_n(iy)} \right| = 1 \quad (y \in \mathbb{R}), \quad \left| \frac{P_n(z)}{Q_n(z)} \right| \leq 1 \quad (\operatorname{Re} z \geq 0), \quad (7.9)$$

the second by the maximum principle, as all roots of Q_n lie in the left half-plane. The error of the approximation is entirely in the phase.

(ii) Bounded phase. Writing $P_n(iy)/Q_n(iy) = e^{i\phi_n(y)}$, $\phi_n = -2 \arg Q_n(iy)$ decreases monotonically and

$$\phi_n(y) = -y + O(y^{2n+1}) \quad (y \rightarrow 0), \quad \phi_n(y) \searrow -n\pi \quad (y \rightarrow \infty), \quad (7.10)$$

whereas the exact phase $-y$ is unbounded (Figure 7.5). This behavior depends only on $y = \omega T$.

(iii) **Right-half-plane zeros.** The n zeros of $P_n(sT)$ have positive real part. The step response of the approximation jumps to $(-1)^n$ at $t = 0^+$ and oscillates before $t = T$, while the true delay stays at zero (Figure 7.6).




Figure 7.5: (a) Unwrapped phase of e^{-iy} and of $[n/n]$ approximations; dotted lines mark $-n\pi$. (b) Phase error; the dashed line marks 10^{-2} rad.




Figure 7.6: Step response of e^{-s} and of its Padé approximations.

Table 7.4 quantifies property (ii): the largest y for which the phase error is below a tolerance. For 0.01 rad and $n \geq 4$, $y_{0.01} \approx 1.6n - 2$.

Table 7.4: Largest $y = \omega T$ with Padé phase error below 0.1, 0.01 and 10^{-6} rad, and asymptotic phase.

n	$y_{0.1}$	$y_{0.01}$	$y_{10^{-6}}$	limit phase $-n\pi$
1	1.124	0.499	0.022	-3.14
2	2.561	1.528	0.235	-6.28
3	4.149	2.813	0.723	-9.42
4	5.819	4.240	1.442	-12.57
6	9.294	7.335	3.369	-18.85
8	12.877	10.621	5.720	-25.13
10	16.526	14.022	8.341	-31.42
12	20.218	17.502	11.149	-37.70

7.8 What Padé preserves

Proposition 7.6 (Frequencies and directions are inherited). *Assume simple crossings, $G \neq 0$ and no common factor of N and D on the imaginary axis. For the rational equation $D + NP_n/Q_n = 0$:*

- (a) *imaginary roots $i\omega$, $\omega > 0$, can occur only at the candidate frequencies ω_i of (7.4) (whether one occurs for a finite delay is decided by (b));*
- (b) *at ω_i , the critical delays solve $\phi_n(\omega_i T) = -(\theta_i + 2\pi k)$, which has a (unique) solution if and only if $\theta_i + 2\pi k < n\pi$;*
- (c) *every crossing has the direction sign $G'(\omega_i)$.*

Diagonal Padé never invents crossing frequencies nor reverses a direction; but a model of order n has at most $\lceil (n\pi - \theta_i)/(2\pi) \rceil$ crossings per frequency for *all* $T \geq 0$, about $n/2$, where the true system has infinitely many. Conclusions that depend only on the modulus are preserved: for P2, $|s + 2| \geq 2 > 1 \geq |P_n/Q_n|$ in the right half-plane, and every Padé model is stable for every delay, as the true system.

In a fixed condition the dominant roots converge very fast with n (Table 7.5, Figure 7.7): faster than geometrically, consistent with the $O(y^{2n+1})$ error, until round-off. Far from the origin, however, rational poles appear that have no exact counterpart (two for $n = 4$, four for $n = 8$, six for $n = 12$ in the window of Figure 7.2).

Table 7.5: RMS error of the dominant pair at the nominal delay.

Case	$n = 1$	$n = 2$	$n = 4$	$n = 6$	$n = 8$	$n = 10$	$n = 12$
P1	5.71e-01	2.90e-02	1.42e-05	1.04e-09	2.29e-14	3.54e-15	1.63e-15
P2	6.53e-01	5.00e-02	4.08e-05	4.73e-09	1.54e-13	3.25e-15	4.09e-15
P3	1.25e-01	5.57e-03	1.20e-06	3.83e-11	1.06e-15	3.58e-15	4.08e-15




Figure 7.7: RMS error of the dominant pair versus Padé order.

7.9 What Padé loses

7.9.1 A non-conservative delay margin

For P1, the $[1/1]$ phase is $-2 \arctan(\omega T/2)$; equating it to $-2\pi/3$ at $\omega_c = \sqrt{3}$ gives $T_c^{[1/1]} = 2$, against the exact 1.2092: a 65 % error (Table 7.6). The rational polynomial $\frac{T}{2}s^2 + (1 - \frac{T}{2})s + 3$ is Hurwitz if and only if $T < 2$. The error *overestimates* the tolerable delay: at $T = 1.5$, the $[1/1]$ model is stable (abscissa $-1/6$) while the true system is unstable (abscissa $+0.0656$). Figure 7.8 confirms it in the time domain, with the delay equation integrated directly by `dde23` [17]. The first crossing occurs at $y = 2\pi/3 \approx 2.09$, beyond the useful range of orders 1 and 2 and within that of order 4 (Table 7.4).




Figure 7.8: P1 at $T = 1.5$, above the exact critical delay 1.209: the delay equation diverges; the $[1/1]$ model decays; $[4/4]$ coincides with the exact solution.

Remark. The rational model also needs an initial state for the Padé filter. The true initial condition is a function on $[-T, 0]$; a model of order n holds only n numbers. The constant history of Figure 7.8 can be represented consistently; general histories cannot.

7.9.2 Shifted and missing branches

For larger delays the branches $k \geq 1$ of (7.5) matter, with larger values of y . Figure 7.9 shows the mechanism for P1: every intersection of a phase curve with a horizontal line $-(\theta + 2\pi k)$ is a critical delay; the Padé phases saturate and reach only the first lines, later and later. Table 7.7 and Figures 7.10–7.11 separate two notions of quality: *local accuracy* of the first critical delay converges fast with n ; *global coverage* of the events grows slowly. In P1 there are three events up to $T = 10$; $[1/1]$ and $[2/2]$ represent one, $[4/4]$ two and $[8/8]$ three. Even when an event is represented, its position may be far off: $[4/4]$ places the second crossing of P1 at 5.71 instead of 4.84.




Figure 7.9: P1: exact and Padé phases versus delay; intersections with the branch lines are critical delays.

Table 7.6: First critical delay: exact value and Padé prediction.

Case and order	predicted T_c	absolute error	relative error
P1 exact	1.209199576	—	—
Padé 1/1	2.000000000	7.908e-01	6.540e-01
Padé 2/2	1.236067977	2.687e-02	2.222e-02
Padé 4/4	1.209215058	1.548e-05	1.280e-05
Padé 8/8	1.209199576	3.575e-14	2.956e-14
Padé 12/12	1.209199576	1.332e-15	1.102e-15
P3 exact	0.291826867	—	—
Padé 1/1	0.305471089	1.364e-02	4.675e-02
Padé 2/2	0.291937845	1.110e-04	3.803e-04
Padé 4/4	0.291826868	9.008e-10	3.087e-09
Padé 8/8	0.291826867	5.551e-17	1.902e-16
Padé 12/12	0.291826867	1.055e-15	3.614e-15

Table 7.7: Coverage of the exact events up to $T = 10$.

Order	events P1	coverage P1	events P3	coverage P3
1	1/3	33.3%	1/6	16.7%
2	1/3	33.3%	2/6	33.3%
3	1/3	33.3%	3/6	50.0%
4	2/3	66.7%	3/6	50.0%
5	2/3	66.7%	4/6	66.7%
6	2/3	66.7%	5/6	83.3%
8	3/3	100.0%	5/6	83.3%
10	3/3	100.0%	6/6	100.0%
12	3/3	100.0%	6/6	100.0%




Figure 7.10: Left: relative error in the first critical delay. Right: fraction of the exact crossings in $[0, 10]$ represented by each order.




Figure 7.11: Critical delays predicted by each order (dots) and exact (vertical lines and crosses).

7.9.3 Stability windows erased, shifted and invented

In P3 stability depends on the order of events at two frequencies, so a shifted event changes the qualitative conclusion. Figure 7.12 compares $Z(T)$ from (7.7) with the rational models (their counts computed with Proposition 7.6 and checked by eigenvalues of a balanced state-space realization), and Table 7.8 lists the predicted stable intervals.




Figure 7.12: P3: number of right-half-plane roots. Black: exact (7.7); circles: independent bound-based count (Proposition 7.2); red: Padé model. Shaded: the rational model reaches the wrong stability conclusion.

Table 7.8: P3: delay intervals in $[0, 10]$ where each model predicts stability, and the number of unstable roots at $T = 10$.

Model	stable intervals in $[0, 10]$	$Z(10)$
exact	$[0, 0.2918), [2.6953, 2.8076)$	4
Padé $[1/1]$	$[0, 0.3055)$	2
Padé $[2/2]$	$[0, 0.2919), (2.8872, 10]$	0
Padé $[3/3]$	$[0, 0.2918), [2.7059, 4.0358)$	2
Padé $[4/4]$	$[0, 0.2918), [2.6957, 2.9797)$	2
Padé $[5/5]$	$[0, 0.2918), [2.6953, 2.8317), (9.1311, 10]$	0
Padé $[6/6]$	$[0, 0.2918), [2.6953, 2.8101)$	2
Padé $[8/8]$	$[0, 0.2918), [2.6953, 2.8076)$	2
Padé $[10/10]$	$[0, 0.2918), [2.6953, 2.8076)$	4
Padé $[12/12]$	$[0, 0.2918), [2.6953, 2.8076)$	4

Four kinds of error appear:

- **Erased window.** $[1/1]$ has no stabilizing event of ω_2 in $[0, 10]$ and misses the window.
- **Unbounded fictitious stability.** $[2/2]$ has the stabilizing event at $T = 2.887$ but not the second branch of ω_1 (which needs $y \approx 7.0 > 2\pi$); it predicts stability for *every* $T > 2.887$, whereas the system is unstable for every $T > 2.808$.

- **Shifted and widened window.** $[3/3]$ and $[4/4]$ predict windows 12 and 2.5 times wider than the true one.
- **Invented stable region.** $[5/5]$ reproduces the window but, lacking the branches $k = 2, 3$ of ω_1 , turns the stabilizing event at $T = 9.13$ into a stable region $(9.13, 10]$ that does not exist.

Only from order 10 on is the stability map in $[0, 10]$ correct, and any fixed order fails again on a long enough interval. High orders are also numerically fragile: the roots of the expanded polynomial $DQ_n + NP_n$ gave a wrong unstable-root count for 24 of 99 small delays with $n = 10$, and 64 of 99 with $n = 12$; a balanced realization avoids this.

7.10 Guidelines

Using Padé safely

1. Decide stability on the original equation: frequencies from (7.4), delays from (7.5), the map (7.7), and a bound-based count as a check.
2. Use Padé as a local tool: design near an operating point, low-order simulation models, or Newton seeds.
3. Choose the order from the largest relevant ωT with Table 7.4; an order n represents no phase beyond $n\pi$.
4. Distrust stability windows and “stable for all larger delays” predicted by rational models.
5. Validate in the time domain near critical delays.
6. Avoid high orders through expanded polynomials; for larger problems use methods designed for delay equations [4, 5, 18].

7.11 Limitations

The direct method above assumes a single constant delay, real coefficients, simple crossings and a retarded system. Tangential or multiple crossings require higher-order analysis [6, 14]. For neutral systems ($p = m$) Proposition 7.2 fails, the root chains approach vertical lines, and stability can be sensitive to infinitesimal changes of the delay [10, 14]. With several delays, crossings form curves and surfaces [13, 15]; matrix systems lead to nonlinear eigenvalue problems [5, 14].

Chapter 8

Distributed-parameter systems

Transfer functions of systems governed by partial differential equations are ratios of entire functions that often look multivalued. This chapter shows how to write them as analytic pairs and computes their poles for the heat, string, duct and beam examples of Curtain and Morris [7]. The models are implemented in `examples/models/distributed_model.m`; the study is `examples/distributed_systems/run_nonrational_examples.m`.

8.1 Removing apparent branch points

For the heat equation the natural expressions contain $\sqrt{s}$. Define

$$C(q) = \cosh \sqrt{q} = \sum_{k \geq 0} \frac{q^k}{(2k)!}, \quad H(q) = \frac{\sinh \sqrt{q}}{\sqrt{q}} = \sum_{k \geq 0} \frac{q^k}{(2k+1)!}, \quad H(0) = 1.$$

Both are entire: only integer powers of q appear. The implementation evaluates H by its series for $|q| < 10^{-4}$ and by the closed form elsewhere, which avoids the $0/0$ at $q = 0$ and does not introduce an artificial pole. The same idea applies to the models below: each factor is evaluated through a regularized formula and is analytic in the search regions used here, so `AssumeAnalytic` is justified for both factors. The heat, string and duct factors are entire; the Kelvin–Voigt beam factors are analytic except at $s = -1/c_d$ (Section 8.4), which the regions must exclude.

8.2 Heat equation: the sensor position

With the normalizations $L = \alpha = K_0 = 1$ of the paper, boundary control and a point sensor at x_0 , the Neumann, Dirichlet and mixed problems give

$$G_N = \frac{C(x_0^2 s)}{s H(s)}, \quad G_D = \frac{x_0 H(x_0^2 s)}{H(s)}, \quad G_M = \frac{x_0 H(x_0^2 s)}{C(s)}.$$

With $x_0 = 1/2$ and $-100 < \operatorname{Re} s < 2$, the Neumann model keeps only 0 and $-4\pi^2$: the candidates $-\pi^2$ and $-9\pi^2$ cancel. The Dirichlet model keeps $-\pi^2$ and $-9\pi^2$ and cancels $-4\pi^2$. The mixed model keeps $-\pi^2/4$, $-9\pi^2/4$ and $-25\pi^2/4$. With the Dirichlet model and $x_0 = 1$, $G \equiv 1$ and every candidate cancels. These are modes of the PDE that the sensor cannot see (Figure 8.1).




Figure 8.1: Poles retained for four sensor positions and three boundary conditions (the Dirichlet row $x_0 = 1$ is empty because $G = 1$). The last panel shows the characteristic example of Chapter 2.

8.3 String and duct

For the string of Section 2 of the paper, with actuator distribution $b(x) = 1$ on $[0, 1/2]$ and $L = c = 1$, the modal coefficients are proportional to $1 - \cos(k\pi/2)$, which vanishes for k multiple of 4; those modes are not poles. The damped model is the output feedback $G/(1 + \varepsilon G)$ of the paper, with $\varepsilon = 0.5$, and the removable factor at $s = 0$ is regularized by a series.

For the acoustic duct with constant reflectance $a = 0.5$ and $L = c = 1$, the poles lie on the vertical line $\ln(a)/2 + ik\pi$. With the radiation impedance of Section 3.4 (Fig. 5 parameters of the paper: $L = 3.54$ m, $c = 341$ m/s, $\rho = 1.20$ kg/m³, radius 0.101 m, sensor at $L/2$), the reflectance depends on frequency and the poles bend to the left. Its rational denominator is cleared from N and D together, so that no artificial poles are introduced (Figure 8.2).

Visible poles of the damped string in the left half-plane do not prove internal stability of modes hidden from this input/output pair.




Figure 8.2: String (undamped, damped, and cancelled modes), duct with constant reflectance and with radiation impedance, and beam poles compared with an independent modal formula.

8.4 Kelvin–Voigt beam: an accumulation point

For the shear-force to tip-velocity transfer function of Section 4.2 of the paper, with $L = EI = I = 1$ and Kelvin–Voigt damping $c_d = 0.02$, let

$$q(s) = \frac{-s^2}{1 + c_d s}, \quad m = q^{1/4}, \quad G(s) = \frac{s [\cosh m \sin m - \sinh m \cos m]/m^3}{(1 + c_d s)[1 + \cosh m \cos m]}.$$

The factors containing m are invariant under $m \mapsto im$, so the choice of fourth root disappears except at $s = -1/c_d$; the numerator ratio has the removable value $2/3$ at $m = 0$. With the positive roots β_k of $1 + \cosh \beta \cos \beta = 0$, the poles solve $s^2 + c_d \beta_k^4 s + \beta_k^4 = 0$, an independent reference. In $[-25, 2] \times [-160, 160]$ the search returns $-0.12362363 \pm 3.51384128i$ and $-4.85518819 \pm 21.49292828i$, within 10^{-14} of the reference.

As $\beta_k \rightarrow \infty$, one branch of these poles converges to $-1/c_d = -50$ (Figure 8.3). Every neighborhood of -50 contains infinitely many poles: G is not meromorphic there, and no finite contour computation can resolve a region containing that point. The model reports it in `meta.singularities`, and the option `Singularities` rejects any region that contains it.




Figure 8.3: Modal reference with 80 modes: one branch of poles converges to $s = -1/c_d = -50$, the other moves to the left.

8.5 Results

Table 8.1: Poles in the stated rectangles; “cancelled” counts locations. The reference errors compare with closed-form or modal formulas that are never used as seeds.

Model	$[\text{Re}_{\min}, \text{Re}_{\max}]$	$ \text{Im} <$	Poles	Cancelled	Max. error
Heat, Neumann	$[-100, 2]$	3	2	2	5×10^{-14}
Heat, Dirichlet	$[-100, 2]$	3	2	1	1×10^{-14}
Heat, mixed	$[-100, 2]$	3	3	0	7×10^{-15}
String, undamped	$[-2, 1]$	30	14	4	3×10^{-21}
String, $\varepsilon = 0.5$	$[-2, 1]$	30	14	4	—
Duct, constant reflectance	$[-2, 1]$	20	13	0	2×10^{-12}
Duct, radiation impedance	$[-400, 30]$	2500	16	0	—
Beam, Kelvin–Voigt	$[-25, 2]$	160	4	0	7×10^{-15}

For the damped string and the radiating duct there is no independent closed form; they are verified by the contour counts and the residuals of the original equations.

Chapter 9

A beam coupled to an oscillator

This case study couples a partial differential equation (a beam, with infinitely many modes) to an ordinary one (a mass-spring-damper), derives an entire characteristic function and transfer functions, and computes the poles without modal truncation. The results are validated with an independent finite-element model.




Figure 9.1: The two topologies. The connection transmits transverse force but no moment; displacements are positive downwards.

9.1 Model

A uniform Euler–Bernoulli beam of length L , bending stiffness EI and mass per length $\mu_b = \rho A$ is clamped at $x = 0$. In topology (a), a mass m with displacement $z(t)$ is connected to the tip displacement $y(t) = w(L, t)$ by a spring k and a damper c in parallel; the mass is not grounded (a tuned-mass absorber). The beam is elastic; all dissipation is in c . Shear deformation, rotary inertia and gravity (already balanced) are neglected:

$$\mu_b w_{tt} + EI w_{xxxx} = 0, \quad 0 < x < L, \quad w(0, t) = w_x(0, t) = 0, \quad (9.1)$$

$$EI w_{xx}(L, t) = 0, \quad -EI w_{xxx}(L, t) = F(t) + f(t), \quad (9.2)$$

$$f = k(z - y) + c(\dot{z} - \dot{y}), \quad m\ddot{z} + c(\dot{z} - \dot{y}) + k(z - y) = u(t), \quad (9.3)$$

where f is the force of the connection on the beam and F, u are external forces at the tip and on the mass.

9.2 The bare-beam receptance

With $w = W(x)e^{st}$ and $\theta^4 = q = -\mu_b L^4 s^2 / EI$, the clamped solution is $W(x) = a[\cosh bx - \cos bx] + d[\sinh bx - \sin bx]$, $b = \theta/L$. The zero-moment condition fixes d/a , and the shear condition gives the tip receptance

$$H_b(s) = \frac{y}{f} = \frac{L^3}{EI} \frac{A(q)}{B(q)}, \quad A(q) = \frac{\cosh \theta \sin \theta - \sinh \theta \cos \theta}{\theta^3}, \quad B(q) = 1 + \cosh \theta \cos \theta.$$

A and B are entire in q : they are invariant under $\theta \mapsto i\theta$. Near $q = 0$, $A(q) = \frac{2}{3} - \frac{q}{315} + \frac{q^2}{623700} - \frac{q^3}{5108103000} + \dots$ and $B(0) = 2$, so $H_b(0) = L^3/(3EI)$, the static compliance of a cantilever. Write $H_b = N_b/D_b$ with $N_b = (L^3/EI)A$, $D_b = B$.

9.3 Characteristic function and transfer functions

With $Z(s) = k + cs$ the Laplace-domain equations are

$$\begin{bmatrix} D_b + N_b Z & -N_b Z \\ -Z & ms^2 + Z \end{bmatrix} \begin{bmatrix} y \\ z \end{bmatrix} = \begin{bmatrix} N_b F \\ u \end{bmatrix},$$

whose determinant is the entire characteristic function

$$\Delta(s) = D_b(s)(ms^2 + cs + k) + N_b(s)ms^2(cs + k). \quad (9.4)$$

The coupled poles are not the union of the beam and oscillator poles, and $1 + H_b Z_{\text{eff}}$, with $Z_{\text{eff}} = ms^2(cs + k)/(ms^2 + cs + k)$, is not entire; clearing its denominators gives (9.4). The transfer functions are

$$\frac{y}{F} = \frac{N_b(ms^2 + Z)}{\Delta}, \quad \frac{y}{u} = \frac{N_b Z}{\Delta}, \quad \frac{z}{u} = \frac{D_b + N_b Z}{\Delta}.$$

With zero inputs, the energy $E = \frac{1}{2} \int_0^L (\mu_b w_t^2 + EI w_{xx}^2) dx + \frac{1}{2} m \dot{z}^2 + \frac{1}{2} k (z - y)^2$ satisfies $\dot{E} = -c(\dot{z} - \dot{y})^2 \leq 0$, which confirms the sign conventions and excludes growing modes for positive parameters. In topology (b) the mass moves with the tip and $\Delta_{\text{tip}} = D_b + N_b(ms^2 + cs + k)$; it is a different system, with potential energy $ky^2/2$.

9.4 Implementation

With $\tau = \sqrt{\mu_b L^4 / EI}$, $\lambda = \tau s$, $\mu = m/(\mu_b L)$, $\kappa = kL^3/EI$ and $\gamma = cL^3/(EI\tau)$, the function evaluated is $\Delta = B(-\lambda^2)(\mu\lambda^2 + \gamma\lambda + \kappa) + A(-\lambda^2)\mu\lambda^2(\gamma\lambda + \kappa)$, equal to (9.4) times L^3/EI . In normalized units the complete computation needs only function handles:

```
m = 0.1; k = 1.2362363368; c = 0.05; % L = EI = rho*A = 1
beta = @(s) (-s.^2).^(1/4);
Nb = @beam_numerator; % series near 0, see text
Db = @(s) 1 + cosh(beta(s)).*cos(beta(s));
Zc = @(s) k + c*s;
N = @(s) Nb(s).*(m*s.^2 + Zc(s));
D = @(s) Db(s).*(m*s.^2 + Zc(s)) + Nb(s).*m*s.^2.*Zc(s);
[p, info] = cpoles(ndpair(N, D), [-60 1 -150 150], 'AssumeAnalytic', true);
```

This excerpt is not self-contained: the local function `beam_numerator` (the series-regularized $A(q)$) is defined in the full script `examples/coupled_beam/beam_direct_nd.m`; the helper `examples/models/coupled_beam_model.m` builds the dimensional model and the three transfer functions.

9.5 Reference case and parameter studies

The reference case uses $m = 0.1$, $k = m\omega_{b1}^2 = 1.23623634$ with the first bare-beam frequency $\omega_{b1} = 3.51601527$, and $c = 0.05$. The ten poles in $[-60, 1] \times [-150, 150]$ are (upper half-plane):

Re(s)	Im(s)	ζ
-0.04655950	2.57427003	0.018084
-0.29481789	4.76361120	0.061771
-0.10690219	22.14678239	0.004827
-0.10105104	61.73606590	0.001637
-0.10032435	120.92154746	0.000830

For $c = 0$ the two lowest frequencies are 2.5703 and 4.7801: the coupling splits the original resonance; in the first mode the mass follows the tip, in the second it moves against it (Figure 9.2).




Figure 9.2: Conservative coupled modes from the exact spatial solution, normalized by the tip displacement.

One parameter at a time was varied: $m \in [0.02, 0.5]$ (25 values), $k \in [0.05, 5]$ (36 values) and $c \in [0, 1.5]$ (43 values). Every one of the 104 cases has ten poles in the window, each found with a fresh contour count (previous poles were only seeds).




Figure 9.3: Low-frequency poles as the mass, the stiffness and the damping vary.

c	$\text{Re}(s)$	$\text{Im}(s)$	ζ
0.00	-0.000000	2.570269	0.0000
0.00	0.000000	4.780068	-0.0000
0.05	-0.046560	2.574270	0.0181
0.05	-0.294818	4.763611	0.0618
0.15	-0.134523	2.606584	0.0515
0.15	-0.892649	4.629622	0.1893
0.30	-0.224003	2.710806	0.0824
0.30	-1.850766	4.152388	0.4071
0.60	-0.192357	2.884138	0.0665
0.60	-4.116381	1.428956	0.9447
1.00	-1.429161	0.000000	1.0000
1.00	-0.123947	2.938350	0.0421
1.00	-13.808315	-0.000000	1.0000
1.50	-0.873874	0.000000	1.0000
1.50	-0.083993	2.954900	0.0284
1.50	-24.605705	0.000000	1.0000

Increasing c can make the second pair real. Among the sampled values the rightmost pole is leftmost at $c = 0.375$ (abscissa -0.234798); at $c = 1.5$ the dominant pair returns to $-0.0840 \pm 2.9549i$. Damping in the connection restrains the very relative motion it dissipates, so more damping is not always better. This is a statement about the sampled values and the stated window, not a global optimum.




Figure 9.4: Lowest branches versus mass and stiffness, window abscissa versus damping, and convergence of the finite-element model.

9.6 Independent validation

A cubic-Hermite finite-element model with consistent mass matrix and an extra degree of freedom for the absorber mass (`coupled_beam_fem.m`) is solved as a quadratic eigenvalue problem,

without using any result of the nonrational solver. The largest error among the ten reference poles decreases with the mesh:

Elements	Max. absolute error
8	0.270955
16	0.0182861
32	0.00116392
64	7.30496e-05

Further checks: the static compliance $H_b(0) = 1/3$; invariance of the dimensionless characteristic function under dimensional rescaling; the bare-beam limit; and the decoupled limit $k = c = 0$, where the characteristic function has a double root at 0 (the free mass) that is *not* a pole of the tip transfer function y/F — a hidden mode. Enlarging the window to $[-90, 1] \times [-300, 300]$ at three damping values leaves the rightmost pole unchanged. The beam is elastic and has no finite accumulation point, unlike the Kelvin–Voigt beam of Section 8.4.

Chapter 10

Validation and limitations

10.1 Test suite

The test suite (`buildtool test` and `buildtool docs`) has four layers.

Unit tests of the public interface: agreement of `croots`, `cpoles`, `czeros` and `cpzmap` with the core solver; polynomial input against `roots`; cancellations; warnings for exploratory and incomplete searches; input errors.

Regression tests of the numerical behavior: a triple root; three roots under scalings of $10^{\pm 100}$; a function without zeros; boundary roots and exhausted budgets that must not be reported as complete; heat, string and duct poles against closed forms; total and partial cancellation; an opaque handle with $Z - P = 0$; symbolic and LTI input; the delay tools (closed-form critical delay, delay-independent stability, Padé predictions, bound-based counts); the coupled beam (static compliance, limits, scaling, finite-element convergence).

Baseline results of the research code from which the toolbox was extracted must be reproduced within a relative tolerance of 10^{-9} .

Documentation tests execute every MATLAB block of the README and of the Markdown documentation, the quick-start scripts, and check that code shown in tutorials is identical to the example files.

10.2 Independent references in the studies

The studies of Chapters 7–9 compare the solver with closed forms (critical delay of P1; heat, string and duct spectra), modal formulas (Kelvin–Voigt beam), the direct critical-delay method and bound-based counts (stability maps), time-domain integration with `dde23`, and a finite-element model (coupled beam). The references are never used as Newton seeds.

10.3 Limitations

- Scalar functions only; matrix and MIMO problems are not treated.
- Rectangular regions only; roots outside the rectangle are not considered, and infinite spectra are not enumerated.
- Completeness is a numerical check under the analyticity assumption (Section 4.7), not an interval-arithmetic proof.
- Analyticity of function handles, and the location of branch and accumulation points, must be established by the user.

- Roots closer than the confirmation resolution are reported as clusters; near pole-zero pairs may be reported as cancelled.
- The time-delay tools assume one constant delay, real coefficients and (for the bound) a retarded system.
- Tested with MATLAB R2023b; other releases and Octave have not been tested.

Appendix A

Proofs

Proposition 7.2. Let $\operatorname{Re} s \geq 0$ and $F(s, T) = 0$. Then $|D(s)| = |N(s)|e^{-T \operatorname{Re} s} \leq |N(s)|$. With $r = |s|$, the triangle inequality gives $|D(s)| \geq r^m - \sum_{k < m} |d_k| r^k$ and $|N(s)| \leq \sum_{k \leq p} |n_k| r^k$, hence $h(r) = r^m - \sum_{k < m} |d_k| r^k - \sum_{k \leq p} |n_k| r^k \leq 0$. Since $p < m$, the coefficients of h have exactly one sign change, so by Descartes' rule of signs h has exactly one positive root R (if some coefficient other than the leading one is nonzero), and $h(r) > 0$ for $r > R$. Therefore $r \leq R$. Nothing depends on T . $\square$

Proposition 7.5. The signs of $\operatorname{Re}(ds/dT)$ and of $\operatorname{Re}[(ds/dT)^{-1}]$ coincide. From (7.6), using $e^{-sT}N(s) = -D(s)$ at a root,

$$\left(\frac{ds}{dT}\right)^{-1} = \frac{D'(s) + e^{-sT}N'(s)}{s e^{-sT}N(s)} - \frac{T}{s} = \frac{1}{s} \left[\frac{N'(s)}{N(s)} - \frac{D'(s)}{D(s)} \right] - \frac{T}{s}.$$

At $s = i\omega$, $-T/(i\omega)$ is purely imaginary and $\operatorname{Re}[\zeta/(i\omega)] = \operatorname{Im}(\zeta)/\omega$. For a real polynomial P , $\frac{d}{d\omega} \ln |P(i\omega)|^2 = 2 \operatorname{Re}[iP'(i\omega)/P(i\omega)] = -2 \operatorname{Im}[P'(i\omega)/P(i\omega)]$. Hence

$$\operatorname{Re}\left(\frac{ds}{dT}\right)^{-1} = \frac{1}{2\omega} \frac{d}{d\omega} [\ln |D(i\omega)|^2 - \ln |N(i\omega)|^2] = \frac{G'(\omega)}{2\omega |D(i\omega)|^2},$$

using $|D| = |N|$ at the crossing. Since $\omega > 0$ the claim follows. $\square$

Proposition 7.6. (a) On the imaginary axis $|P_n/Q_n| = 1$ by (7.9), so $D + NP_n/Q_n = 0$ requires $|D(i\omega)| = |N(i\omega)|$. (b) The phase condition is $\phi_n(\omega T) = -\theta_i - 2\pi k$; since ϕ_n decreases monotonically from 0 to $-n\pi$, it has a solution, which is unique, if and only if $\theta_i + 2\pi k < n\pi$. (c) With $R(z) = P_n(z)/Q_n(z)$ in place of e^{-z} , $F = D + R(sT)N$, $F_s = D' + RN' + TR'N$ and $F_T = sR'N$. Using $RN = -D$ and $\rho = R'/R$,

$$\left(\frac{ds}{dT}\right)^{-1} = -\frac{F_s}{F_T} = \frac{1}{s[-\rho(sT)]} \left[\frac{N'}{N} - \frac{D'}{D} \right] - \frac{T}{s}.$$

On the imaginary axis $R(iy) = e^{i\phi_n(y)}$ gives $\rho(iy) = \phi'_n(y)$, real and negative; for the exact delay $\rho \equiv -1$. The positive factor $-1/\phi'_n$ does not change the sign of the real part, which is therefore $\operatorname{sign} G'(\omega_i)$ as in Proposition 7.5. $\square$

Error of the derivative formula of Section 4.2. For analytic f , $\frac{f(z+h)-f(z-h)}{2h} = f' + \frac{h^2}{6} f''' + O(h^4)$ and $\frac{f(z+ih)-f(z-ih)}{2ih} = f' - \frac{h^2}{6} f''' + O(h^4)$. Their mean is the formula of Section 4.2, with error $O(h^4)$. $\square$

Appendix B

Reproducibility

All numbers, tables and figures of this manual are produced by the code distributed with ContourRoots. From the repository folder:

```
buildtool examples % time-delay, distributed and coupled-beam studies
buildtool manual   % documentation figures and this PDF
buildtool test     % unit and regression tests
buildtool docs     % documentation tests
```

File	Content
matlab/complex_spectrum.m	core solver (Chapter 4)
matlab/private/spectrum_*.m	counts, Newton, subdivision, input adapters
matlab/croots.m, cpoles.m, czeros.m, cpzmap.m, ndpair.m	public interface
matlab/delay/*.m	time-delay and Padé tools (Chapter 7)
examples/time_delay/ run_delay_study.m	Chapter 7; calls <code>run_pade_critical_study</code> and <code>run_pade_pitfalls</code>
examples/distributed_systems/ run_nonrational_examples.m	Chapter 8
examples/coupled_beam/ run_coupled_beam_study.m	Chapter 9
examples/models/*.m	distributed and coupled-beam models, finite elements
tools/make_doc_figures.m	figures of the Markdown documentation
manual/	this manual (L ^A T _E X sources and TikZ figures)

The results in this version were computed with MATLAB R2023b. Generated data are written to `output/`, which is not versioned; the studies take a few minutes.

Authorship and AI assistance

ContourRoots was created and is maintained by Flávio Luiz Cardoso-Ribeiro. It was developed extensively with AI coding assistance, using Claude Code (Anthropic) and Codex (OpenAI), including code generation, documentation and tests; Codex was also used for an independent review of version 0.1.0. Problem formulation, project direction and release decisions are led by the maintainer. Numerical validation and known limitations are documented in Chapter 10, in the Markdown documentation (`docs/validation.md`, `docs/diagnostics_and_limits.md`) and in `ROADMAP.md`.

Bibliography

- [1] Lars V. Ahlfors. *Complex Analysis*. McGraw-Hill, New York, 3 edition, 1979.
- [2] George A. Baker and Peter Graves-Morris. *Padé Approximants*, volume 59 of *Encyclopedia of Mathematics and its Applications*. Cambridge University Press, 2 edition, 1996.
- [3] Richard Bellman and Kenneth L. Cooke. *Differential-Difference Equations*. Academic Press, New York, 1963.
- [4] Dimitri Breda, Stefano Maset, and Rossana Vermiglio. Pseudospectral differencing methods for characteristic roots of delay differential equations. *SIAM Journal on Scientific Computing*, 27(2):482–495, 2005. doi: 10.1137/030601600.
- [5] Dimitri Breda, Stefano Maset, and Rossana Vermiglio. *Stability of Linear Delay Differential Equations: A Numerical Approach with MATLAB*. Springer, 2015. doi: 10.1007/978-1-4939-2107-2.
- [6] Kenneth L. Cooke and P. van den Driessche. On zeroes of some transcendental equations. *Funkcialaj Ekvacioj*, 29:77–90, 1986.
- [7] Ruth F. Curtain and Kirsten Morris. Transfer functions of distributed parameter systems: A tutorial. *Automatica*, 45(5):1101–1116, 2009. doi: 10.1016/j.automatica.2009.01.008.
- [8] L. M. Delves and J. N. Lyness. A numerical method for locating the zeros of an analytic function. *Mathematics of Computation*, 21(100):543–560, 1967.
- [9] Keqin Gu, Vladimir L. Kharitonov, and Jie Chen. *Stability of Time-Delay Systems*. Birkhäuser, 2003. doi: 10.1007/978-1-4612-0039-0.
- [10] Jack K. Hale and Sjoerd M. Verduyn Lunel. *Introduction to Functional Differential Equations*, volume 99 of *Applied Mathematical Sciences*. Springer, 1993. doi: 10.1007/978-1-4612-4342-7.
- [11] Piotr Kowalczyk. GRPF: Global complex roots and poles finding algorithm. <https://github.com/PioKow/GRPF>. Accessed 2026-09-26.
- [12] Peter Kravanja and Marc Van Barel. *Computing the Zeros of Analytic Functions*, volume 1727 of *Lecture Notes in Mathematics*. Springer, Berlin, 2000.
- [13] Wim Michiels and Silviu-Iulian Niculescu. Characterization of delay-independent stability and delay interference phenomena. *SIAM Journal on Control and Optimization*, 45(6): 2138–2155, 2007. doi: 10.1137/050641223.
- [14] Wim Michiels and Silviu-Iulian Niculescu. *Stability, Control, and Computation for Time-Delay Systems: An Eigenvalue-Based Approach*. SIAM, 2 edition, 2014. doi: 10.1137/1.9781611973631.
- [15] Nejat Olgac and Rifat Sipahi. An exact method for the stability analysis of time-delayed linear time-invariant (LTI) systems. *IEEE Transactions on Automatic Control*, 47(5): 793–797, 2002. doi: 10.1109/TAC.2002.1000275.

- [16] Robert Parini. cxroots: A Python module to find all the roots of a complex analytic function within a given contour. <https://github.com/rparini/cxroots>. Accessed 2026-09-26.
- [17] L. F. Shampine and S. Thompson. Solving DDEs in MATLAB. *Applied Numerical Mathematics*, 37(4):441–458, 2001. doi: 10.1016/S0168-9274(00)00055-6.
- [18] Tomáš Vyhlídal and Pavel Zítek. Mapping based algorithm for large-scale computation of quasi-polynomial zeros. *IEEE Transactions on Automatic Control*, 54(1):171–177, 2009.
- [19] K. Walton and J. E. Marshall. Direct method for TDS stability analysis. *IEE Proceedings D – Control Theory and Applications*, 134(2):101–107, 1987.